

PROJECT TITLE : RECOGNIZING HANDWRITTEN DIGITS USING MACHINE LEARNING

PROGRAM:

```
# Import required libraries

import numPy as np

import tensorflow as tf

from tensorflow.keras.datasets import mnist

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

from tensorflow.keras.utils import to_categorical


# Enable memory growth for GPU (optional)

physical_devices = tf.config.list_physical_devices('GPU')

if physical_devices:

    try:

        tf.config.set_memory_growth(physical_devices[0], True) # Updated API

    except RuntimeError as e:

        print(f"GPU memory growth setting failed: {e}")


# Load the MNIST dataset

(X_train, y_train), (X_test, y_test) = mnist.load_data()


# Reshape and normalize the input data

X_train = X_train.reshape(-1, 28, 28, 1).astype('float32') / 255.0

X_test = X_test.reshape(-1, 28, 28, 1).astype('float32') / 255.0


# Convert class labels to one-hot encoding
```

```
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

# Build the CNN model
model = Sequential([
    Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D(pool_size=(2, 2)),

    Conv2D(64, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),

    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(10, activation='softmax') # 10 classes for digits 0–9
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Print model summary
model.summary()

# Train the model
model.fit(X_train, y_train, batch_size=128, epochs=5, validation_split=0.1)

# Evaluate the model
loss, accuracy = model.evaluate(X_test, y_test)
print("\nTest Accuracy: {:.4f}".format(accuracy))
```

output:

```
Model: "sequential_1"
...


| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d_2 (Conv2D)              | (None, 26, 26, 32) | 320     |
| max_pooling2d_2 (MaxPooling2D) | (None, 13, 13, 32) | 0       |
| conv2d_3 (Conv2D)              | (None, 11, 11, 64) | 18,496  |
| max_pooling2d_3 (MaxPooling2D) | (None, 5, 5, 64)   | 0       |
| flatten_1 (Flatten)            | (None, 1600)       | 0       |
| dense_2 (Dense)                | (None, 128)        | 204,928 |
| dropout_1 (Dropout)            | (None, 128)        | 0       |
| dense_3 (Dense)                | (None, 10)         | 1,290   |


Total params: 225,034 (879.04 KB)
Trainable params: 225,034 (879.04 KB)
Non-trainable params: 0 (0.00 B)
Epoch 1/5
422/422 ————— 41s 93ms/step - accuracy: 0.7987 - loss: 0.6566 - val_accuracy: 0.9815 - val_loss: 0.0670
Epoch 2/5
422/422 ————— 39s 94ms/step - accuracy: 0.9674 - loss: 0.1141 - val_accuracy: 0.9858 - val_loss: 0.0521
Epoch 3/5
422/422 ————— 42s 95ms/step - accuracy: 0.9762 - loss: 0.0794 - val_accuracy: 0.9875 - val_loss: 0.0428
Epoch 4/5
422/422 ————— 39s 92ms/step - accuracy: 0.9811 - loss: 0.0616 - val_accuracy: 0.9897 - val_loss: 0.0386
Epoch 5/5
401/422 ————— 1s 88ms/step - accuracy: 0.9833 - loss: 0.0551
```